

Factors of Success in Indie Games: A Data-Driven Analysis of Steam Reviews

Semester Project Report

Zhamshid Kabanbay

Data Collection & Analysis

Platform: Steam (store.steampowered.com)

Dataset: 1,486 indie games | 110,748 reviews

2025

1. Introduction

The independent game development market has experienced remarkable growth over the past decade. As of 2024, Steam — the world's largest PC gaming platform — hosts over 50,000 titles, with indie games accounting for more than 70% of all releases. Despite this volume, the vast majority of indie games fail to achieve commercial or critical success.

Understanding what separates successful indie games from unsuccessful ones is a question of both academic and practical importance. Developers, publishers, and investors lack systematic, data-driven insights into which game attributes — genre, mechanics, technical quality, narrative depth, or pricing — are most strongly associated with player satisfaction.

This study addresses this gap by analyzing 110,748 user reviews from 1,486 indie games collected from Steam. Using a combination of thematic text analysis, sentiment analysis, tag-based lift analysis, and multivariate regression, we identify statistically significant factors associated with game success, as measured by the Wilson score lower bound.

Research Questions:

- RQ1: Which thematic dimensions of user reviews (bugs, narrative, gameplay, price, developer support, etc.) differentiate successful from unsuccessful indie games?
- RQ2: What is the relationship between the overall sentiment of user reviews and game success?
- RQ3: Which genre tags are most associated with high-rated indie games?

2. Materials and Methods

2.1 Data Collection

Data were collected using a custom Python scraper that interfaced with two public APIs: the Steam Store API (store.steampowered.com/api) and the SteamSpy API (steamspy.com/api.php). The collection pipeline proceeded in three stages:

- Stage 1 — Game Discovery: Games tagged as 'Indie' were retrieved from SteamSpy using the genre endpoint. Only games with a minimum of 10 user reviews were included to eliminate near-zero-activity titles.
- Stage 2 — Filtering: Each candidate game was verified against the Steam Store API to confirm it was classified as type 'game' (excluding DLC, soundtracks, and software).

- Stage 3 — Review Collection: Up to 100 English-language reviews per game were collected via the Steam Reviews API using cursor-based pagination.

The final dataset comprised 110,748 reviews across 1,486 unique indie games. A checkpoint system ensured fault-tolerant collection, saving progress after each game.

2.2 Success Metric: Wilson Score

Game success was operationalized using the Wilson score lower bound, a statistically robust metric that accounts for both the proportion of positive reviews and the total number of reviews:

$$\text{score} = (p + z^2/2n - z\sqrt{((p(1-p) + z^2/4n)/n)}) / (1 + z^2/n)$$

where p is the proportion of positive reviews, n is the total number of reviews, and $z = 1.96$ (95% confidence interval). This metric is preferred over simple positive ratios because it penalizes games with very few reviews, preventing a game with 3 out of 3 positive reviews from outranking a game with 9,000 out of 10,000.

2.3 Sample Stratification

To enable valid group comparisons, games were divided into two equal groups by median Wilson score (median = 0.767):

- Successful games: Wilson score ≥ 0.767 ($n = 744$ games, 61,046 reviews)
- Unsuccessful games: Wilson score < 0.767 ($n = 742$ games, 49,702 reviews)

Median-based stratification was chosen over arbitrary thresholds to ensure balanced group sizes, which is a prerequisite for unbiased statistical comparison.

2.4 Thematic Analysis

Eight thematic dimensions were defined a priori based on common factors discussed in game development literature and player feedback frameworks:

Theme	Example Keywords
Bugs & Technical Issues	bug, crash, broken, glitch, freeze, lag
Narrative & Story	story, character, plot, writing, dialogue, ending
Gameplay & Mechanics	gameplay, controls, combat, satisfying, challenge

Theme	Example Keywords
Price & Value	price, refund, worth, expensive, money, cheap
Graphics & Art	graphics, art, visual, style, beautiful, pixel
Music & Sound	music, soundtrack, sound, audio, ost
Content & Length	short, long, replay, content, finished, hours
Developer Support	dev, update, abandoned, patch, community

For each review, a binary indicator was computed per theme (1 if any keyword matched, 0 otherwise). Theme frequency was then calculated as the proportion of reviews mentioning the theme within each group.

2.5 Statistical Tests

Two statistical measures were applied to each theme comparison:

- Chi-square test of independence (χ^2): Tests whether theme mention frequency differs significantly between successful and unsuccessful games. A p-value < 0.05 indicates statistically significant difference.
- Cohen's d (effect size): Quantifies the practical magnitude of the difference. Interpreted as: d < 0.2 = negligible, d = 0.2–0.5 = small, d = 0.5–0.8 = medium, d ≥ 0.8 = large.

Both measures are reported together, as p-values alone can be misleading with large samples — statistically significant effects may be practically negligible.

2.6 Sentiment Analysis

Review-level sentiment was computed using VADER (Valence Aware Dictionary and sEntiment Reasoner), a lexicon-based sentiment analyzer optimized for short social media texts. Each review received a compound sentiment score ranging from –1 (maximally negative) to +1 (maximally positive). Game-level sentiment was computed as the mean compound score across all reviews for a given game. Pearson correlation was used to assess the association between mean sentiment and Wilson score.

2.7 Tag Lift Analysis

To identify genre tags associated with highly-rated games, a lift metric was computed for each tag:

$$\text{lift}(\text{tag}) = \text{frequency}(\text{tag} \mid \text{top 25\%}) / \text{frequency}(\text{tag} \mid \text{all games})$$

A lift > 1.0 indicates the tag appears more frequently among top-quartile games; lift < 1.0 indicates it is underrepresented. Only tags appearing in at least 10 games were included.

2.8 Regression Analysis

A multivariate Ordinary Least Squares (OLS) regression model was estimated with Wilson score as the dependent variable and the eight thematic frequencies plus mean sentiment as independent variables. The model was estimated using the statsmodels library in Python. Cross-validation (5-fold) R^2 was computed using scikit-learn to assess generalizability.

3. Results and Discussion

3.1 Thematic Analysis Results

Table 2 presents the thematic analysis results with chi-square statistics, p-values, and Cohen's d effect sizes.

Theme	Successful %	Unsuccessful %	Diff %	χ^2	p-value	Cohen's d	Effect
Bugs & Technical Issues	6.9%	14.9%	-8.0%	1894.1	< 0.001	0.261	Small
Developer Support	11.9%	18.8%	-6.9%	1034.9	< 0.001	0.193	Negligible
Price & Value	17.7%	22.8%	-5.1%	450.4	< 0.001	0.128	Negligible
Narrative & Story	23.7%	21.2%	+2.5%	101.5	< 0.001	0.061	Negligible
Gameplay & Mechanics	19.4%	21.0%	-1.6%	47.2	< 0.001	0.041	Negligible

Theme	Successful %	Unsuccessful %	Diff %	χ^2	p-value	Cohen's d	Effect
Music & Sound	20.0%	21.4%	-1.4%	32.8	< 0.001	0.035	Negligible
Graphics & Art	27.1%	28.1%	-1.0%	13.3	0.0003	0.022	Negligible
Content & Length	24.4%	25.0%	-0.6%	5.5	0.019	0.014	Negligible




Figure 1. Thematic analysis results: frequency comparison (left) and effect sizes (right). * = $p < 0.05$.

All eight themes showed statistically significant differences between groups ($p < 0.05$). However, only Bugs & Technical Issues reached a small effect size ($d = 0.261$), while all other themes showed negligible effects. This distinction is critical: statistical significance here is partly an artifact of the large sample size ($n = 110,748$ reviews), which provides sufficient power to detect even trivially small differences.

The most substantively meaningful finding is that bug-related discourse was more than twice as prevalent in unsuccessful game reviews (14.9% vs. 6.9%), suggesting that poor technical quality is a primary driver of player dissatisfaction. Developer support discourse followed a similar pattern (18.8% vs. 11.9%), with terms such as 'abandoned' showing a particularly strong within-theme effect ($d = 0.136$).

Narrative & Story was the only theme more prevalent in successful game reviews (+2.5%), suggesting that games with strong narrative elements tend to receive more positive evaluations. However, the effect size ($d = 0.061$) indicates this association is weak.

3.2 Sentiment Analysis

The Pearson correlation between game-level mean sentiment and Wilson score was $r = 0.538$ ($p < 0.001$), indicating a moderate positive association. Games with more positively-toned reviews tend to receive higher Wilson scores, which is conceptually consistent — players who rate a game positively also tend to write positively about it.

Mean sentiment was substantially higher for successful games compared to unsuccessful games, confirming that sentiment polarity effectively discriminates between the two groups at the aggregate level.

3.3 Tag Lift Analysis




Figure 2. Left: Top-20 tags by mean Wilson score. Right: Tag regression coefficients ($R^2 = 0.097$).

The lift analysis revealed that niche, community-oriented tags are most strongly associated with high-rated indie games. The top tags by lift include:

- Word Game (lift = 2.21), Co-op Campaign (2.12), Underwater (2.06), Parkour (2.04), Cozy (1.67), LGBTQ+ (1.68), Precision Platformer (1.68)

These tags share a common characteristic: they represent specific, well-defined player niches with loyal audiences whose expectations align precisely with the game's offering. A 'Cozy' game player knows what to expect; a 'Word Game' player has a specific demand being met.

Conversely, tags with the lowest lift (anti-tags) include:

- Battle Royale (lift = 0.27), 4X (0.24), Modern (0.24), Combat Racing (0.35), Open World (0.49), City Builder (0.43)

These genres represent highly competitive, AAA-dominated spaces where indie developers struggle to meet player expectations set by major studio titles. An indie Battle Royale, for instance, is inevitably compared to Fortnite or PUBG — a comparison few indie teams can survive.

3.4 Regression Analysis




Figure 3. OLS regression coefficients. * = $p < 0.05$. Blue = positive effect, Red = negative effect.

The OLS regression model achieved $R^2 = 0.502$ and Adjusted $R^2 = 0.498$, indicating that the nine predictors (eight thematic frequencies + mean sentiment) jointly explain approximately 50% of variance in Wilson score. Cross-validation confirmed generalizability (5-fold CV $R^2 \approx 0.49 \pm 0.03$).

Mean sentiment was by far the strongest predictor ($\beta \approx 0.52$, $p < 0.001$), consistent with its high bivariate correlation with Wilson score. Among thematic predictors, Narrative & Story showed a positive coefficient, while Price & Value, Bugs & Technical Issues, and Developer Support showed negative coefficients — all statistically significant.

Notably, Gameplay & Mechanics showed a small positive coefficient in the regression despite being slightly more prevalent in unsuccessful game reviews in the bivariate analysis. This sign reversal illustrates the value of multivariate analysis: when controlling for bugs and developer support, gameplay-focused discourse is associated with slightly higher scores, suggesting these reviews reflect engaged players rather than disgruntled ones.

3.5 Summary of Key Findings

- Technical quality (bugs, crashes) is the strongest thematic predictor of poor ratings ($d = 0.261$, $\beta < 0$), reinforcing that baseline technical competence is a prerequisite for indie success.
- Developer abandonment is a strong signal of failure: 'abandoned' appeared in 3.4% of unsuccessful game reviews vs. 0.4% of successful ones.
- Niche positioning outperforms broad appeal: games in well-defined niches (Word Game, Cozy, Precision Platformer) consistently outperform games competing in saturated mainstream genres (Battle Royale, 4X, Open World).
- Sentiment explains ~54% of variance in Wilson score ($r = 0.538$), making it the strongest single predictor in the dataset.
- No single thematic factor dominates: the plurality of negligible effect sizes suggests that indie game success is multifactorial and cannot be reduced to any single dimension.

4. Conclusion

This study analyzed 110,748 Steam user reviews across 1,486 indie games to identify statistically significant factors associated with game success, measured by the Wilson score lower bound. Using thematic analysis, sentiment analysis, tag lift analysis, and multivariate OLS regression, we identified several actionable patterns.

The most robust finding is that technical quality — specifically the absence of bugs and the presence of ongoing developer support — is the single most important differentiator between successful and unsuccessful indie games. Games where players frequently discuss bugs, crashes, and abandonment consistently receive lower Wilson scores, with effect sizes reaching $d = 0.261$ for bug-related discourse, the only finding to exceed the 'negligible' threshold.

A secondary finding concerns genre positioning: indie games in niche, well-defined genres (Cozy, Word Game, Precision Platformer, LGBTQ+) show lift values above 2.0, while games

in AAA-competitive genres (Battle Royale, 4X, Open World) show lift values below 0.5. This suggests that strategic genre selection — choosing niches underserved by major studios rather than competing head-on — is a viable path to relative success for indie developers.

The regression model ($R^2 = 0.502$) demonstrates that sentiment and thematic content of user reviews are moderately predictive of game success, accounting for half of the variance in Wilson score. The remaining variance likely reflects factors outside the scope of this analysis, including marketing reach, release timing, pricing strategy, and random viral effects.

Limitations: This study is limited to English-language Steam reviews, which may introduce linguistic and cultural bias. The keyword-based thematic approach may miss nuanced expressions not captured by the predefined lexicon. The sample is restricted to indie games and may not generalize to other game categories. Future work should incorporate multilingual reviews, temporal analysis of review dynamics, and pricing data as an additional predictor.

5. References

Note: This report is based on primary data collected directly from Steam and SteamSpy public APIs. The following tools and libraries were used in the analysis:

- Steam Store API. Valve Corporation.
<https://store.steampowered.com/api/appdetails>
- SteamSpy API. SteamSpy. <https://steamspy.com/api.php>
- Hutto, C., & Gilbert, E. (2014). VADER: A Parsimonious Rule-Based Model for Sentiment Analysis of Social Media Text. *Proceedings of the International AAAI Conference on Web and Social Media*, 8(1).
- Wilson, E. B. (1927). Probable inference, the law of succession, and statistical inference. *Journal of the American Statistical Association*, 22(158), 209–212.
- Cohen, J. (1988). *Statistical Power Analysis for the Behavioral Sciences* (2nd ed.). Lawrence Erlbaum Associates.
- Pedregosa, F., et al. (2011). Scikit-learn: Machine Learning in Python. *Journal of Machine Learning Research*, 12, 2825–2830.
- McKinney, W. (2010). Data Structures for Statistical Computing in Python. *Proceedings of the 9th Python in Science Conference*, 51–56.

- Seabold, S., & Perktold, J. (2010). Statsmodels: Econometric and Statistical Modeling with Python. Proceedings of the 9th Python in Science Conference.